

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

Question 1

(a)

(a) Naïve Bayes defines the joint probability of the each data point \mathbf{x} and its class label c as follows: $p(\mathbf{x}, c | \boldsymbol{\theta}, \boldsymbol{\pi}) = p(c | \boldsymbol{\theta}, \boldsymbol{\pi}) p(\mathbf{x} | c, \boldsymbol{\theta}, \boldsymbol{\pi}) = p(c | \boldsymbol{\pi}) \prod_{j=1}^{784} p(x_j | c, \theta_{jc})$, and knowing that $p(x_j | c, \theta_{jc}) = \theta_{jc}^{x_j} (1 - \theta_{jc})^{(1-x_j)}$, and $p(c | \boldsymbol{\pi}) = \pi_c$, we have:

$L(\boldsymbol{\theta}) = p(\mathbf{x}, c | \boldsymbol{\theta}, \boldsymbol{\pi}) = \prod_{i=1}^N \pi_c^i \prod_{j=1}^{784} \theta_{jc}^{x_j^i} (1 - \theta_{jc})^{(1-x_j^i)}$, in this way we could derive log-likelihood $l(\boldsymbol{\theta})$:

$$l(\boldsymbol{\theta}) = \frac{\partial L(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \sum_{i=1}^N \left(\log \pi_c^i + \sum_{j=1}^{784} (x_j^i \log \theta_{jc} + (1 - x_j^i) \log (1 - \theta_{jc})) \right)$$

Thus, $\hat{\boldsymbol{\theta}}_{MLE} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} l(\boldsymbol{\theta})$ s.t. with respect to $\sum_{c=0}^9 \pi_c = 1$.

Take derivative with respect to θ_{jc} , we have $\frac{\partial l(\boldsymbol{\theta})}{\partial \theta_{jc}} = \sum_{i=1}^N \mathbb{1}(C^{(i)} = c) \left(\frac{x_j^i}{\theta_{jc}} - \frac{1-x_j^i}{1-\theta_{jc}} \right) = 0$

$$\implies \sum_{i=1}^N \mathbb{1}(C^{(i)} = c) (x_j^i (1 - \theta_{jc}) - \theta_{jc} (1 - x_j^i)) = 0$$

$$\implies \sum_{i=1}^N \mathbb{1}(C^{(i)} = c) \theta_{jc} = \sum_{i=1}^N \mathbb{1}(C^{(i)} = c) x_j^i$$

$$\implies \hat{\theta}_{MLE} = \frac{\sum_{i=1}^N \mathbb{1}(C^{(i)} = c \wedge x_j^i = 1)}{\sum_{i=1}^N \mathbb{1}(C^{(i)} = c)}.$$

Now derive MLE for $\boldsymbol{\pi}$. By using Lagrange multiplier, we have

$$\frac{\partial l(\boldsymbol{\theta})}{\partial \pi_j} + \lambda \frac{\partial \sum \pi_j}{\partial \pi_j} = 0 \implies \lambda = - \sum_{n=1}^N I(t^{(n)} = j) \frac{1}{\pi_j} \implies \hat{\pi}_j = \frac{- \sum_{n=1}^N I(t^{(n)} = j)}{\lambda}.$$

Apply constraint $\sum_{c=0}^9 \pi_c = 1$, we have $\lambda = -N$. In this way, we have

$$\hat{\pi}_j = \frac{\sum_{n=1}^N I(t^{(n)} = j)}{N}$$

(b) (c)

$$\begin{aligned} \text{(b)} \quad p(t | \mathbf{x}, \boldsymbol{\theta}, \boldsymbol{\pi}) &= \frac{p(\mathbf{x}, c | \boldsymbol{\theta}, \boldsymbol{\pi})}{p(\mathbf{x} | \boldsymbol{\theta}, \boldsymbol{\pi})} = \frac{p(\mathbf{x}, c | \boldsymbol{\theta}, \boldsymbol{\pi})}{\sum_{c=0}^9 p(\mathbf{x}, c | \boldsymbol{\theta}, \boldsymbol{\pi})} = \frac{\pi_c \prod_{j=1}^{784} \theta_{jc}^{x_j} (1 - \theta_{jc})^{1-x_j}}{\sum_{c=0}^9 \pi_c \prod_{j=1}^{784} \theta_{jc}^{x_j} (1 - \theta_{jc})^{1-x_j}}. \\ \implies \log(p(t | \mathbf{x}, \boldsymbol{\theta}, \boldsymbol{\pi})) &= \log(\pi_c \prod_{j=1}^{784} (\theta_{jc}^{x_j} (1 - \theta_{jc})^{1-x_j})) - \log(\sum_{c=0}^9 \pi_c \prod_{j=1}^{784} (\theta_{jc}^{x_j} (1 - \theta_{jc})^{1-x_j})) \\ &= \log(\pi_c) + \sum_{j=1}^{784} x_j \log(\theta_{jc}) + (1 - x_j) \log(1 - \theta_{jc}) \\ &\quad - \sum_{c=0}^9 \left(\log(\pi_c) + \sum_{j=1}^{784} x_j \log(\theta_{jc}) + (1 - x_j) \log(1 - \theta_{jc}) \right) \\ &= - \sum_{i=0, i \neq c}^9 \left(\log(\pi_i) + \sum_{j=1}^{784} (x_j) \log(\theta_{ij}) + \log(1 - \theta_{ij}) (1 - x_j) \right) \end{aligned}$$

Step 1 load data

```
1 from __future__ import absolute_import
2 from __future__ import print_function
3 from future.standard_library import install_aliases
4 install_aliases()
5 import numpy as np
6 import os
7 import gzip
8 import struct
9 import array
10 import matplotlib.pyplot as plt
11 import matplotlib.image
12 from urllib.request import urlretrieve
13 from scipy.special import logsumexp
14 from IPython.display import display, Math, Latex
```

```
1 def download(url, filename):
2     if not os.path.exists('data'):
```

```

2  if not os.path.exists('data'):
3      os.makedirs('data')
4  out_file = os.path.join('data', filename)
5  if not os.path.isfile(out_file):
6      urlretrieve(url, out_file)
7
8
9  def mnist():
10     base_url = 'http://yann.lecun.com/exdb/mnist/'
11
12     def parse_labels(filename):
13         with gzip.open(filename, 'rb') as fh:
14             magic, num_data = struct.unpack(">II", fh.read(8))
15             return np.array(array.array("B", fh.read()), dtype=np.uint8)
16
17     def parse_images(filename):
18         with gzip.open(filename, 'rb') as fh:
19             magic, num_data, rows, cols = struct.unpack(">IIII", fh.read(16))
20             return np.array(array.array("B", fh.read()), dtype=np.uint8).reshape(num_data, rows, cols)
21
22     for filename in ['train-images-idx3-ubyte.gz',
23                     'train-labels-idx1-ubyte.gz',
24                     't10k-images-idx3-ubyte.gz',
25                     't10k-labels-idx1-ubyte.gz']:
26         download(base_url + filename, filename)
27
28     train_images = parse_images('data/train-images-idx3-ubyte.gz')
29     train_labels = parse_labels('data/train-labels-idx1-ubyte.gz')
30     test_images = parse_images('data/t10k-images-idx3-ubyte.gz')
31     test_labels = parse_labels('data/t10k-labels-idx1-ubyte.gz')
32
33     return train_images, train_labels, test_images[:1000], test_labels[:1000]
34
35
36 def load_mnist():
37     partial_flatten = lambda x: np.reshape(x, (x.shape[0], np.prod(x.shape[1:])))
38     one_hot = lambda x, k: np.array(x[:, None] == np.arange(k)[None, :], dtype=int)
39     train_images, train_labels, test_images, test_labels = mnist()
40     train_images = (partial_flatten(train_images) / 255.0 > .5).astype(float)
41     test_images = (partial_flatten(test_images) / 255.0 > .5).astype(float)
42     train_labels = one_hot(train_labels, 10)
43     test_labels = one_hot(test_labels, 10)
44     N_data = train_images.shape[0]
45
46     return N_data, train_images, train_labels, test_images, test_labels
47
48
49 def plot_images(images, ax, ims_per_row=5, padding=5, digit_dimensions=(28, 28),
50                 cmap=matplotlib.cm.binary, vmin=None, vmax=None):
51     """Images should be a (N_images x pixels) matrix."""
52     N_images = images.shape[0]
53     N_rows = np.int32(np.ceil(float(N_images) / ims_per_row))
54     pad_value = np.min(images.ravel())
55     concat_images = np.full(((digit_dimensions[0] + padding) * N_rows + padding,
56                             (digit_dimensions[1] + padding) * ims_per_row + padding), pad_value)
57     for i in range(N_images):
58         cur_image = np.reshape(images[i, :], digit_dimensions)
59         row_ix = i // ims_per_row
60         col_ix = i % ims_per_row
61         row_start = padding + (padding + digit_dimensions[0]) * row_ix
62         col_start = padding + (padding + digit_dimensions[1]) * col_ix
63         concat_images[row_start: row_start + digit_dimensions[0],
64                       col_start: col_start + digit_dimensions[1]] = cur_image
65     cax = ax.matshow(concat_images, cmap=cmap, vmin=vmin, vmax=vmax)
66     plt.xticks(np.array([]))
67     plt.yticks(np.array([]))
68     return cax
69
70
71 def save_images(images, filename, **kwargs):
72     fig = plt.figure(1)
73     fig.clf()
74     ax = fig.add_subplot(111)
75     plot_images(images, ax, **kwargs)
76     fig.patch.set_visible(False)
77     ax.patch.set_visible(False)
78     plt.savefig(filename)
79
80 if __name__ == '__main__':
81     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
82     print(N_data)
83     print(train_images.shape)
84     print(train_labels.shape)
85     print(test_images.shape)
86     print(test_labels.shape)

```

```

↳ 60000
   (60000, 784)
   (60000, 10)
   (1000, 784)
   (1000, 10)

```

▼ Step 2: According to question a, complete train_mle_estimator function

```

1 def train_mle_estimator(train_images, train_labels):
2     """ Inputs: train_images, train_labels
3         Returns the MLE estimators theta_mle and pi_mle """
4     N, D = train_images.shape
5     theta_mle = np.zeros((10, D))
6     pi_mle = np.zeros(10)
7
8     for i in range(10):
9         pi_mle[i] = sum(train_labels[:, i]) / N
10
11    for c in range(10):
12        for d in range(D):
13            theta_mle[c][d] = (1 + (np.sum(np.where((train_images[:, d] == 1) & (train_labels[:, c] == 1),
14            1, 0)))) / (2 + (np.sum(train_labels[:, c])))
15
16
17    return theta_mle, pi_mle

```

▼ Step3: Write log-likelihood function according to the equation derived

```

1 def log_likelihood_with_loop(images, theta, pi, train_labels):
2     N, D = images.shape
3     log_like = np.zeros((N, 10))
4
5     for i in range(N):
6         for c in range(10):
7             if train_labels[i][c] != 1:
8                 log_like[i][c] -= np.log(pi[c])
9                 for d in range(D):
10                     log_like[i][c] -= images[i][d] * np.log(theta[c][d])
11                     log_like[i][c] -= np.log(1 - theta[c][d]) * (1 - images[i][d])
12    return log_like

```

▼ Problem 1:

since there are so many layers of loop in this function, it would drastically slow down the running speed of the program.

```

1 def log_likelihood(images, theta, pi, train_labels):
2     """ Inputs: images, theta, pi
3         Returns the matrix 'log_like' of loglikelihoods over the input images where
4         log_like[i,c] = log p (c | x^(i), theta, pi) using the estimators theta and pi.
5         log_like is a matrix of num of images x num of classes
6         Note that log likelihood is not only for c^(i), it is for all possible c's. """
7
8     N = images.shape[0]
9     C, D = theta.shape
10
11    log_like = np.zeros((N, C))
12
13    for i in range(N):
14        bern = np.where(images[i] > 0.5, theta, 1- theta)
15        temp = np.log(pi) + np.sum(np.log(bern), axis = 1)
16        log_like[i] = temp - logsumexp(temp)
17    return log_like
18
19 def average_log_likelihood(images, theta, pi, train_labels):
20     log_like = log_likelihood(images, theta, pi, train_labels)
21     average_log_like = np.average(log_like, axis = 0)
22     return average_log_like
23
24 if __name__ == '__main__':
25     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
26     theta_mle, pi_mle = train_mle_estimator(train_images, train_labels)
27     avg_loglike_mle = average_log_likelihood(train_images, theta_mle, pi_mle, train_labels)
28     print(avg_loglike_mle)

```

```

↳ [-109.28662699 -181.68961943 -72.47438253 -72.9793892 -73.12914282
   -55.01175952 -94.88549624 -105.08574599 -59.50921323 -81.3683617 ]

```


Problem 2: Notice that if we don't handle the situation for log0, the system would report "RuntimeWarning: divide by zero encountered in log", and the result would be -inf. In this case, we should make some minor edit on the log-likelihood function to avoid the situation of log0.

▼ (d)

Now we plot the MLE estimator $\hat{\theta}$ as 10 separate greyscale images, one for each class.

```
1 if __name__ == '__main__':
2     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
3     theta_mle, pi_mle = train_mle_estimator(train_images, train_labels)
4
5     fig = plt.figure(1)
6     fig.clf()
7     ax = fig.add_subplot(111)
8     plot_images(theta_mle, ax)
```

↪ <matplotlib.image.AxesImage at 0x7efc7a9f4668>



▼ (e)

(e) Since $P(\theta|\mathbf{x}, c, \pi) \propto P(\theta)P(\mathbf{x}, c|\theta, \pi)$ where $\theta_{jc} \sim \text{Beta}(3, 3)$, we could calculate $l(\theta)$ as:

$$l(\theta) = \log P(\theta) + \log P(\mathbf{x}, c|\theta, \pi) + c, \text{ where } c \text{ is constant}$$

$$= \log(\theta_{jc}(1 - \theta_{jc})) + \log \pi_c^i + \sum_{j=1}^{784} (x_j^i \log \theta_{jc} + (1 - x_j^i) \log(1 - \theta_{jc})) + c.$$

Similar to the analysis in question a, by taking derivative with respect to θ_{jc} we have:

$$\frac{\partial l(\theta)}{\partial \theta_{jc}} = \left(\frac{1}{\theta_{jc}} - \frac{1}{1 - \theta_{jc}} \right) + \sum_{i=1}^N \mathbb{1}(C^i = c) \left(\frac{x_j^i}{\theta_{jc}} - \frac{1 - x_j^i}{1 - \theta_{jc}} \right) = 0$$

$$\implies (1 - \theta_{jc}) - \theta_{jc} + \sum_{i=1}^N \mathbb{1}(C^i = c) (x_j^i (1 - \theta_{jc}) - (1 - x_j^i) \theta_{jc}) = 0$$

$$\implies \sum_{i=1}^N \mathbb{1}(C^i = c) \theta_{jc} + 2\theta_{jc} = \sum_{i=1}^N \mathbb{1}(C^i = c) x_j^i + 1$$

$$\implies \hat{\theta}_{MLE} = \frac{1 + \sum_{i=1}^N \mathbb{1}(C^i = c \wedge x_j^i = 1)}{2 + \sum_{i=1}^N \mathbb{1}(C^i = c)}$$

```
1 def train_map_estimator(train_images, train_labels):
2     """ Inputs: train_images, train_labels
3     Returns the MAP estimators theta_map and pi_map"""
4     N, D = train_images.shape
5     theta_mle = np.zeros((10, D))
6     pi_mle = np.zeros(10)
7
8     for i in range(10):
9         pi_mle[i] = sum(train_labels[:, i]) / N
10
11     for c in range(10):
12         for d in range(D):
13             theta_mle[c][d] = (1 + (np.sum(np.where((train_images[:, d] == 1) & (train_labels[:, c] == 1),
14                 1, 0)))) / (2 + (np.sum(train_labels[:, c])))
15
16     return theta_mle, pi_mle
```

▼ (f)

Similar to question c, derive average log-likelihood per data point and and the accuracy on both the training and test set.

```
1 def predict(log_like):
2     """ Inputs: matrix of log likelihoods
3     Returns the predictions based on log likelihood values"""
4     return np.argmax(log_like, axis=1)
5
6
7 def accuracy(log_like, labels):
8     """ Inputs: matrix of log likelihoods and 1-of-K labels
9     Returns the accuracy based on predictions from log likelihood values"""
10    prediction = predict(log_like)
11    return np.mean(prediction == np.argmax(labels, axis = 1))
12
13 if __name__ == '__main__':
14     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
15     theta_map, pi_map = train_map_estimator(train_images, train_labels)
```

```

15 theta_map, pi_map = train_map_estimator(train_images, train_labels)
16 loglike_train_map = log_likelihood(train_images, theta_map, pi_map, train_labels)
17 avg_loglike_map = np.sum(loglike_train_map * train_labels) / N_data
18 print("Average log-likelihood for MAP is ", avg_loglike_map)
19 print("LOADED TRAIN LABEL ", train_labels.shape)
20 train_accuracy_map = accuracy(loglike_train_map, train_labels)
21 loglike_test_map = log_likelihood(test_images, theta_map, pi_map, test_labels)
22 test_accuracy_map = accuracy(loglike_test_map, test_labels)
23
24 print("Training accuracy for MAP is ", train_accuracy_map)
25 print("Test accuracy for MAP is ", test_accuracy_map)

```


Average log-likelihood for MAP is -3.3558565594023526
 LOADED TRAIN LABEL (60000, 10)
 Training accuracy for MAP is 0.8357666666666667
 Test accuracy for MAP is 0.816

▼ (g)

```

1 if __name__ == '__main__':
2     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
3     theta_map, pi_map = train_map_estimator(train_images, train_labels)
4
5     fig = plt.figure(1)
6     fig.clf()
7     ax = fig.add_subplot(111)
8     plot_images(theta_map, ax)

```



▼ Question 2

(a)

True

Reason:

According to the Naïve Bayes model assumption, this statement is true.

(b)

False

Reason:

Because $p(x_i, x_j) = \sum_c p(x_i, x_j | c) = \sum_c p(x_i | c) p(x_j | c)$ and $p(x_i) p(x_j) = \sum_c p(x_i | c) \sum_c p(x_j | c)$, we could state that $p(x_i, x_j) \neq p(x_i) p(x_j)$. Therefore, the statement that any two pixels x_i and x_j where $i \neq j$ are independent after marginalizing over c is False.

▼ (c)

```

1 def image_sampler(theta, pi, num_images):
2     """ Inputs: parameters theta and pi, and number of images to sample
3     Returns the sampled images"""
4
5     sampled_images = np.ndarray(shape = (10, 784))
6     c = np.random.choice(10, 10, p=pi)
7     print("random c: ", c)
8     for i in range(num_images):
9         sampled_images[i] = np.random.binomial(n=1, p=theta[c[i]])
10    return sampled_images
11
12 if __name__ == '__main__':
13     N_data, train_images, train_labels, test_images, test_labels = load_mnist()
14     theta_map, pi_map = train_map_estimator(train_images, train_labels)
15     sampled_images = image_sampler(theta_map, pi_map, 10)
16

```

```

17 fig = plt.figure(1)
18 fig.clf()
19 ax = fig.add_subplot(111)
20 plot_images(sampled_images, ax)

```

random c: [6 2 3 0 6 4 1 3 9 7]



Question 3

(a)


```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 from sklearn.decomposition import PCA
4 from sklearn.neighbors import KNeighborsClassifier
5 from scipy.spatial import distance

1 def load_data(filename, load2=True, load3=True):
2     """Loads data for 2's and 3's
3     Inputs:
4         filename: Name of the file.
5         load2: If True, load data for 2's.
6         load3: If True, load data for 3's.
7     """
8     assert (load2 or load3), "Atleast one dataset must be loaded."
9     data = np.load(filename)
10    # print(data['train2'].shape)
11
12    if load2 and load3:
13        inputs_train = np.hstack((data['train2'], data['train3']))
14        inputs_valid = np.hstack((data['valid2'], data['valid3']))
15        inputs_test = np.hstack((data['test2'], data['test3']))
16        target_train = np.hstack((np.zeros((1, data['train2'].shape[1])), np.ones((1, data['train3'].shape[1]))))
17        target_valid = np.hstack((np.zeros((1, data['valid2'].shape[1])), np.ones((1, data['valid3'].shape[1]))))
18        target_test = np.hstack((np.zeros((1, data['test2'].shape[1])), np.ones((1, data['test3'].shape[1]))))
19    else:
20        if load2:
21            inputs_train = data['train2']
22            target_train = np.zeros((1, data['train2'].shape[1]))
23            inputs_valid = data['valid2']
24            target_valid = np.zeros((1, data['valid2'].shape[1]))
25            inputs_test = data['test2']
26            target_test = np.zeros((1, data['test2'].shape[1]))
27        else:
28            inputs_train = data['train3']
29            target_train = np.zeros((1, data['train3'].shape[1]))
30            inputs_valid = data['valid3']
31            target_valid = np.zeros((1, data['valid3'].shape[1]))
32            inputs_test = data['test3']
33            target_test = np.zeros((1, data['test3'].shape[1]))
34
35    return inputs_train.T, inputs_valid.T, inputs_test.T, target_train.T, target_valid.T, target_test.T
36
37 if __name__ == '__main__':
38     path = '/content/drive/My Drive/Colab Notebooks/digits.npz'
39     inputs_train, inputs_valid, inputs_test, target_train, target_valid, target_test = load_data(path)
40     a = np.reshape(inputs_train[0], (16, 16))
41     plt.imshow(a, cmap = "gray")
42     plt.show()

```

↳



```


1 inputs_target = np.hstack((np.zeros((300)), np.ones((300))))
2 valid_target = np.hstack((np.zeros((100)), np.ones((100))))
3
4 knn = KNeighborsClassifier(n_neighbors=1)
5
6 inputs_train_mean = np.mean(inputs_train, axis=0)
7 inputs_train_centered = inputs_train - np.tile(inputs_train_mean, (inputs_train.shape[0], 1))
8 inputs_valid_centered = inputs_valid - np.tile(inputs_train_mean, (inputs_valid.shape[0], 1))
9
10 cov_matrix = np.cov(inputs_train_centered.T)
11 eigen_value, eigen_vector = np.linalg.eig(cov_matrix)
12 k_list = [2, 5, 10, 20, 30]
13 mismatch_list = []
14
15 for item in k_list:
16     eigen_matrix = eigen_vector[:, 0:item]
17     inputs_train_recon = np.matmul(inputs_train_centered, eigen_matrix)
18     inputs_valid_recon = np.matmul(inputs_valid_centered, eigen_matrix)
19     knn.fit(inputs_train_recon, inputs_target)
20     inputs_valid_target_pred = knn.predict(inputs_valid_recon)
21
22     counter = 0
23     for i in range(200):
24         if inputs_valid_target_pred[i] != 0 and i <= 99:
25             counter += 1
26         elif inputs_valid_target_pred[i] != 1 and i >= 100:
27             counter += 1
28     mismatch_list.append(counter / 200)
29
30 print("WHEN K VALUE IS ", item, " THE CLASSIFICATION ERROR RATE IS ", counter / 200)
31
32
33 plt.plot(k_list, mismatch_list, label="mismatch_line")
34 plt.legend(loc="upper right")
35 plt.xlabel(r"$k$ value$")
36 plt.ylabel("error rate")
37 plt.show()

```

```

(600, 2)
WHEN K VALUE IS 2 THE CLASSIFICATION ERROR RATE IS 0.055
(600, 5)
WHEN K VALUE IS 5 THE CLASSIFICATION ERROR RATE IS 0.035
(600, 10)
WHEN K VALUE IS 10 THE CLASSIFICATION ERROR RATE IS 0.01
(600, 20)
WHEN K VALUE IS 20 THE CLASSIFICATION ERROR RATE IS 0.005
(600, 30)
WHEN K VALUE IS 30 THE CLASSIFICATION ERROR RATE IS 0.01

```



(b)

I would choose K value at 20. Observe from the data above we notice that, when $K = 20$, the classification error reaches the lowerst point, which is 0.005. When K exceeds 20 the model would overfit, and it would learn some error into the system causing the classification rate

rising. Therefore, I would choose model with K value at 20.

▼ (c)

```


1 test_target = np.hstack((np.zeros((200)), np.ones((200))))
2
3
4 inputs_train_mean = np.mean(inputs_train, axis = 0)
5 inputs_train_centered = inputs_train - np.tile(inputs_train_mean, (inputs_train.shape[0], 1))
6 inputs_test_centered = inputs_test - np.tile(inputs_train_mean, (inputs_test.shape[0], 1))
7 cov_matrix = np.cov(inputs_train_centered.T)
8 eigen_value, eigen_vector = np.linalg.eig(cov_matrix)
9
10 mismatch_list = []
11 for item in k_list:
12     eigen_matrix = eigen_vector[:, 0: item]
13     inputs_train_recon = np.matmul(inputs_train_centered, eigen_matrix)
14     inputs_test_recon = np.matmul(inputs_test_centered, eigen_matrix)
15
16     knn.fit(inputs_train_recon, inputs_target)
17     inputs_test_target_pred = knn.predict(inputs_test_recon)
18
19     counter = 0
20     for i in range(400):
21         if inputs_test_target_pred[i] != test_target[i]:
22             counter += 1
23     mismatch_list.append(counter / 400)
24     print("WHEN K VALUE IS ", item, " THE CLASSIFICATION ERROR RATE IS ", counter / 400)
25
26
27 plt.plot(k_list, mismatch_list, label="mismatch_line")
28 plt.legend(loc="upper right")
29 plt.xlabel(r"$k$ value$")
30 plt.ylabel("error rate")
31 plt.show()

```

```

↳ WHEN K VALUE IS 2 THE CLASSIFICATION ERROR RATE IS 0.065
   WHEN K VALUE IS 5 THE CLASSIFICATION ERROR RATE IS 0.04
   WHEN K VALUE IS 10 THE CLASSIFICATION ERROR RATE IS 0.0125
   WHEN K VALUE IS 20 THE CLASSIFICATION ERROR RATE IS 0.01
   WHEN K VALUE IS 30 THE CLASSIFICATION ERROR RATE IS 0.0075

```



In this way, we choose $K = 30$, where the classification error rate is the lowest, at 0.0075